

xlapes02

December 16, 2023

1 Import Packages

```
[125]: from pathlib import Path
import numpy as np
import matplotlib.pyplot as plt
import scipy.stats as stats
import pandas as pd
from scipy import stats
import statsmodels.api as sm
import statsmodels.formula.api as smf
import statsmodels.stats.api as sms
from statsmodels.stats.outliers_influence import variance_inflation_factor
import seaborn as sns
from sklearn.preprocessing import PolynomialFeatures
from statsmodels.stats.outliers_influence import OLSInfluence
```

2 Setup basic configuration + define helper functions

```
[126]: def write_to_file(file_name, content):
        with open(file_name, 'w') as f:
            f.write(content)

        # Create directory for output files
        Path('tmp/out').mkdir(parents=True, exist_ok=True)

        # Set dark theme
        plt.style.use('dark_background')

        # Set grid thickness
        plt.rcParams['grid.linewidth'] = 0.3
```

3 Load Data

```
[127]: excel_file = pd.ExcelFile("Projekt-2_Data.xlsx")
df_uloha_1: pd.DataFrame = excel_file.parse(excel_file.sheet_names[0])
df_uloha_2 = excel_file.parse(excel_file.sheet_names[1])
data = {
    '1': df_uloha_1,
    '1_a': df_uloha_1['uloha_1 a'],
    '1_b_prior': df_uloha_1['uloha_1 b)_prior'],
    '1_g': df_uloha_1['skupina'],
    '1_b_observation': df_uloha_1['uloha_1 b)_pozorování'],
    '2': df_uloha_2,
    '2_os': df_uloha_2['OSType'],
    '2_as': df_uloha_2['ActiveUsers'],
    '2_ip': df_uloha_2['InteractingPct'],
    '2_sp': df_uloha_2['ScrollingPct'],
    '2_p': df_uloha_2['Ping [ms]'],
}
# data
```

```
/Users/zlapik/.pyenv/versions/3.10.13/lib/python3.10/xml/etree/ElementTree.py:16
47: ResourceWarning: unclosed file <_io.BufferedReader
name='Projekt-2_Data.xlsx'>
    attrib = {}
ResourceWarning: Enable tracemalloc to get the object allocation traceback
```

4 ULOHA 1 - Bayesovske odhady

4.1 ULOHA 1.a - Konjugované apriorní a aposteriorní rozdělení, prediktivní rozdělení [2 body]

4.1.1 Clean data

- Remove outliers
- Remove nan values
- Remove +-inf values
- Remove values with Z-score > 3
- Remove values with Z-score < -3

```
[128]: df_uloha_1 = data['1_a']

# Extract observed data
observed_data = df_uloha_1.values

# Remove nan or +-inf values
observed_data = observed_data[~np.isnan(observed_data)]

# Calculate Z-scores
```

```

z_scores = stats.zscore(observed_data, nan_policy='raise')

# Define a threshold for outliers (e.g., 3 standard deviations)
threshold = 3

# Filter out rows with Z-scores beyond the threshold
filtered_data = observed_data[(np.abs(z_scores) < threshold)]
filtered_data

```

```

[128]: array([2., 2., 1., 3., 0., 1., 1., 3., 2., 2., 3., 1., 5., 3., 1., 1., 2.,
            1., 1., 1., 2., 3., 2., 0., 3., 1., 2., 1., 5., 1., 0., 0., 2., 1.,
            1., 0., 0., 1., 3., 1., 0., 1., 2., 0., 1., 3., 0., 1., 1., 4., 1.,
            2., 1., 1., 2., 4., 2., 2., 3., 4., 4., 4., 0., 2., 0., 0., 3., 5.,
            1., 2., 1., 0., 1., 1., 4., 1., 1., 3., 0., 1., 2., 2., 2., 3., 1.,
            2., 2., 2., 1., 2., 2., 1., 0., 1., 1., 3., 0., 3., 1., 1.])

```

4.2 ULOHA 1.a.1 - Do jednoho obrázku vykreslíte apriorní a aposteriorní hustotou parametru Poissonova rozdělení λ .

```

[129]: alpha_prior = 10 # connection count
       beta_prior = 5 # time within the connection count (alpha_prior) was observed
       lambda_expert = alpha_prior / beta_prior # expert's estimate of the connection_
           ↪ count

```

```

[130]: alpha_posterior = alpha_prior + np.sum(filtered_data)
       beta_posterior = beta_prior + len(filtered_data)

```

```

[131]: x_prior = np.linspace(0, np.max(filtered_data), 1000)
       y_prior = stats.gamma.pdf(x_prior, alpha_prior, scale=1 / beta_prior)

```

```

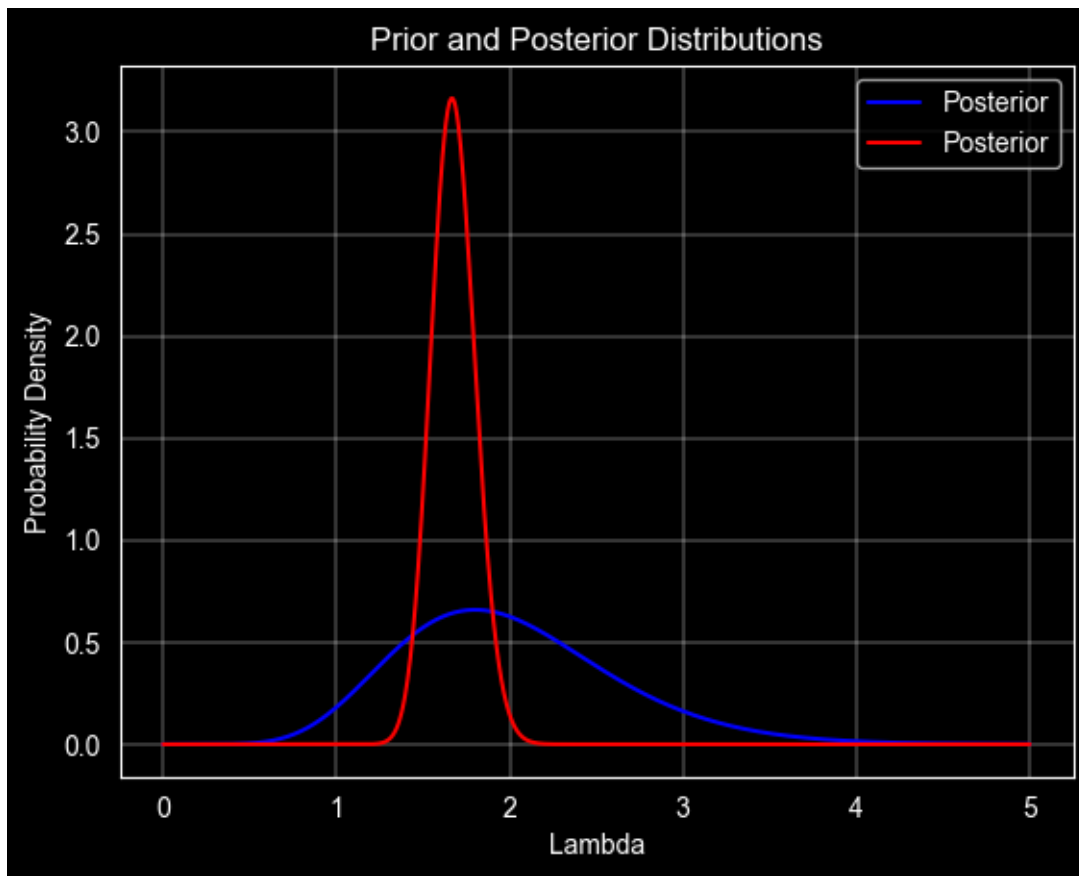
[132]: x_posterior = np.linspace(0, np.max(filtered_data), 1000)
       y_posterior = stats.gamma.pdf(x_posterior, alpha_posterior, scale=1 /
           ↪ beta_posterior)

```

```

[133]: plt.plot(x_prior, y_prior, label='Posterior', color='blue')
       plt.plot(x_posterior, y_posterior, label='Posterior', color='red')
       plt.title('Prior and Posterior Distributions')
       plt.xlabel('Lambda')
       plt.ylabel('Probability Density')
       plt.legend()
       plt.show()

```



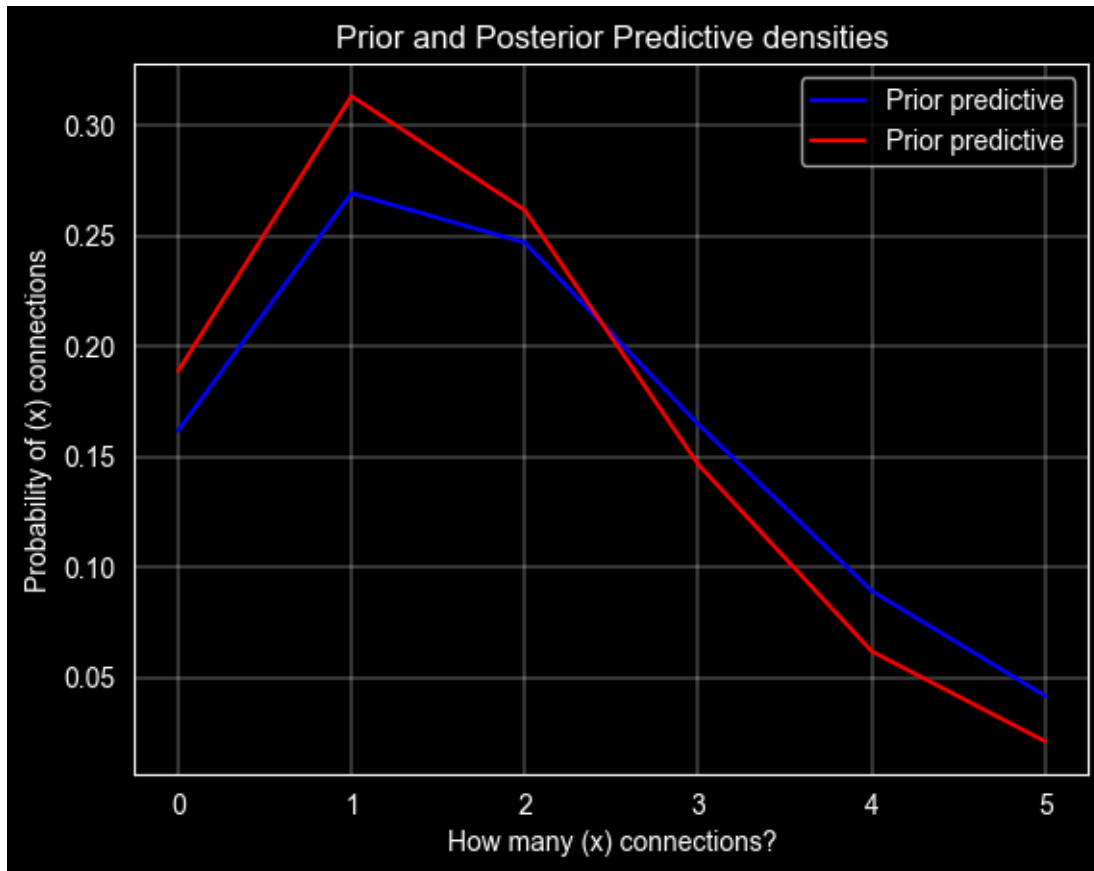
4.3 ULOHA 1.a.2 - Do jednoho obrázku vykreslíte apriorní a aposteriorní prediktivní hustotou pozorování x za jeden časový interval.

```
[134]: x_prior_interval = range(0, 6) # 0 to 5 connections, because nbinom is discrete
y_prior_interval = stats.nbinom.pmf(x_prior_interval, alpha_prior, beta_prior / (1 + beta_prior))
```

```
[135]: x_posterior_interval = range(0, 6) # 0 to 5 connections, because nbinom is discrete
y_posterior_interval = stats.nbinom.pmf(x_posterior_interval, alpha_posterior, beta_posterior / (1 + beta_posterior))
```

```
[136]: plt.plot(x_prior_interval, y_prior_interval, label='Prior predictive', color='blue')
plt.plot(x_posterior_interval, y_posterior_interval, label='Prior predictive', color='red')
plt.title('Prior and Posterior Predictive densities')
plt.xlabel('How many (x) connections?')
plt.ylabel('Probability of (x) connections')
```

```
plt.legend()
plt.show()
```



```
[137]: # Task 3: Construct 95% confidence intervals for from prior and posterior
        ↳ distributions
prior_ci = stats.gamma.interval(0.95, alpha_prior, scale=1 / beta_prior)
posterior_ci = stats.gamma.interval(0.95, alpha_posterior, scale=1 /
        ↳ beta_posterior)
print(f"Prior 95% CI: {prior_ci[0]:.5f}, {prior_ci[1]:.5f}")
print(f"Posterior 95% CI: {posterior_ci[0]:.5f}, {posterior_ci[1]:.5f}")
```

Prior 95% CI: 0.95908, 3.41696

Posterior 95% CI: 1.43769, 1.93272

```
[138]: # Task 4: Select two posterior point estimates for and compare them
posterior_mean = alpha_posterior / beta_posterior
posterior_mode = (alpha_posterior - 1) / beta_posterior
print(f"Aposteriori mean: {posterior_mean:.5f}")
print(f"Aposteriori mode: {posterior_mode:.5f}")
```

Aposteriori mean: 1.67619

Apriori mode: 1.66667

```
[139]: # Task 5: Select one prior and one posterior point estimate for the number of
        ↪ observations
mu_prior = alpha_prior / beta_prior
mu_posterior = alpha_posterior / beta_posterior
print(f"Prior estimate: {mu_prior:.5f}")
print(f"Posterior estimate: {mu_posterior:.5f}")
```

Prior estimate: 2.00000

Posterior estimate: 1.67619

4.4 ULOHA 1.b - Aproximace diskrétním rozdělením [2 body]

```
[140]: mu = 3
sigma = np.sqrt(1)
a = 1
```

4.4.1 Load data

```
[141]: # Cleaned data
df_uloha_1_b = {
    'prior_data': data['1_b_prior'][~np.isnan(data['1_b_prior'])],
    'observed_data': data['1_b_observation'][~np.
    ↪ isnan(data['1_b_observation'])],
    'group_column': data['1_g'][~np.isnan(data['1_g'])]
}
observed_data = df_uloha_1_b['observed_data']
```

4.4.2 Uloha 1.b.1: Plot prior, posterior, and likelihood functions

```
[142]: bins_count = 50

# Get max value for each group
all_data_max = data['1'].groupby('skupina')['uloha_1 b)_prior'].max()

bin_width = (all_data_max.max() - all_data_max.min()) / bins_count # Get bin
    ↪ width
bins = np.arange(all_data_max.min(), all_data_max.max(), bin_width) # Bin values

bin_height, bin_edges = np.histogram(all_data_max, bins=bins_count)
bin_height = bin_height / np.sum(bin_height)

# Plot bins
# plt.bar(x=bins, height=bin_height, width=bin_width, color='blue', alpha=0.7)
# plt.show()
```

```
[143]: bin_centers = (bin_edges[:-1] + bin_edges[1:]) / 2

def likelihood_func(observed_data, b):
    """
    Calculate likelihood function
    :param observed_data:
    :param b:
    :return:
    """
    a_truncnorm = (a - mu) / sigma
    b_truncnorm = (b - mu) / sigma
    pdf = stats.truncnorm.pdf(observed_data, a=a_truncnorm, b=b_truncnorm,
    ↪loc=mu, scale=sigma)
    return pdf

# Calculate likelihood function for each bin
likelihood = [likelihood_func(observed_data, b_center) for b_center in
    ↪bin_centers]

# Calculate product of all likelihoods
likelihood = np.prod(likelihood, axis=1)

# Normalize likelihood
likelihood_normalized = likelihood / np.sum(likelihood)

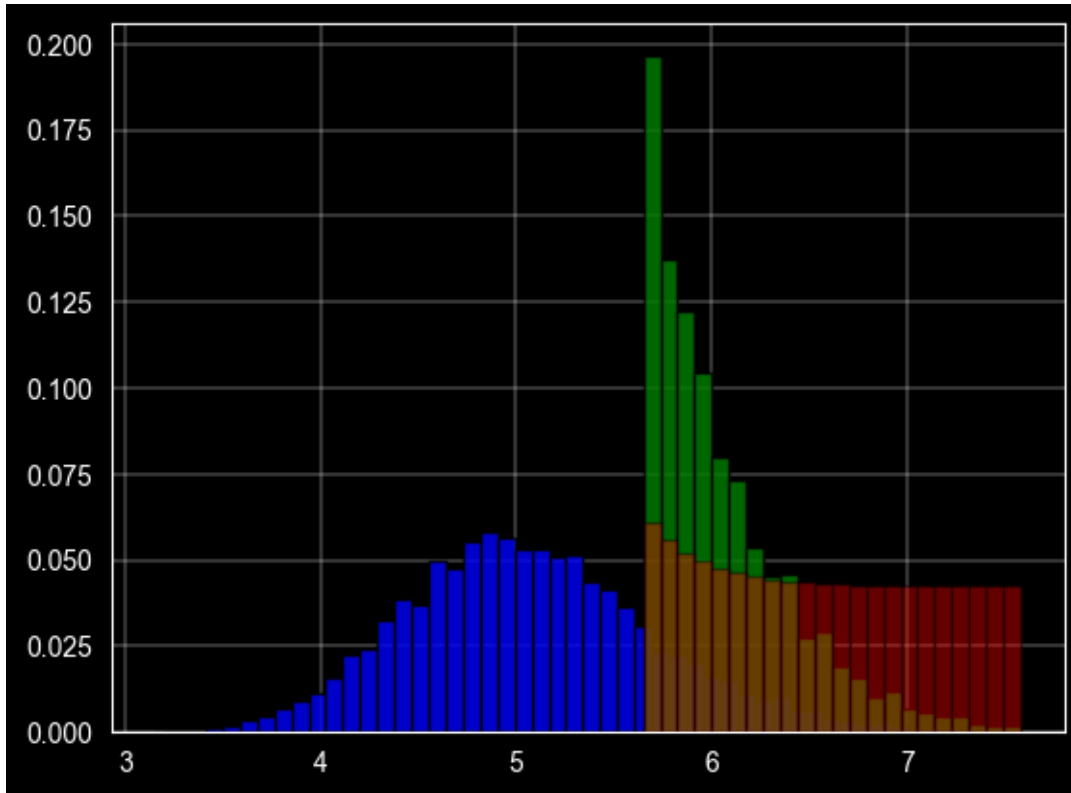
# Plot likelihood
# plt.bar(x=bins, height=likelihood_normalized, width=bin_width,
    ↪edgecolor='black', color='red', label='Likelihood', alpha=0.7)
# plt.show()

[144]: # Calculate posterior
posterior_probs = likelihood * bin_height
posterior_probs_normalized = posterior_probs / np.sum(posterior_probs)

# plt.bar(bin_centers, posterior_probs_normalized, width=bin_width,
    ↪edgecolor='black', color='green', label='Aposteriórne rozdelenie', alpha=0.7)
# plt.show()

[145]: # Plot all together: prior, likelihood, posterior
plt.bar(x=bins, height=bin_height, width=bin_width, color='blue', alpha=0.8,
    ↪edgecolor='black')
plt.bar(bin_centers, posterior_probs_normalized, width=bin_width,
    ↪edgecolor='black', color='green',
    label='Aposteriórne rozdelenie', alpha=0.8)
plt.bar(bin_centers, likelihood_normalized, width=bin_width, edgecolor='black',
    ↪color='red', label='Vierohodnosť',
    alpha=0.4)
```

```
plt.show()
```



4.4.3 Uloha 1.b.2. Z aposteriorní hustoty určete 95% interval spolehlivosti (konfidenční interval) pro parametr .

```
[146]: # Calculate 95% confidence interval
cumulative_posterior = np.cumsum(posterior_probs_normalized)
lower_bound = bin_centers[np.argmax(cumulative_posterior >= 0.025)]
upper_bound = bin_centers[np.argmin(cumulative_posterior <= 0.975)]
print(f'95% Confidence Interval for Parameter b: {lower_bound:.5f},  
↪{upper_bound:.5f}')
```

95% Confidence Interval for Parameter b: 5.69371, 7.00891

4.4.4 Uloha 1.b.3. Vyberte dva bodové odhady parametru a spočítejte je.

```
[147]: # Calculate point estimates
mean = np.sum(bin_centers * posterior_probs_normalized)
median = bin_centers[np.argmax(posterior_probs_normalized)]
print(f'First point estimate: {mean:.5f}')
print(f'Second point estimate: {median:.5f}') # Is this value OK?
```


First point estimate: 6.05277
Second point estimate: 5.69371

5 ULOHA 2 - ÚLOHA 2 – Regrese – 8. bodů

5.1 ULOHA 2.1. [4. body] Pomocí zpětné eliminace určete vhodný regresní model. Za výchozí „plný“ model považujte plný kvadratický model (všechny interakce druhého řádu a všechny druhé mocniny, které dávají smysl).

5.2 Learn more about data

```
[148]: # Load data
df_uloha_2 = data['2']

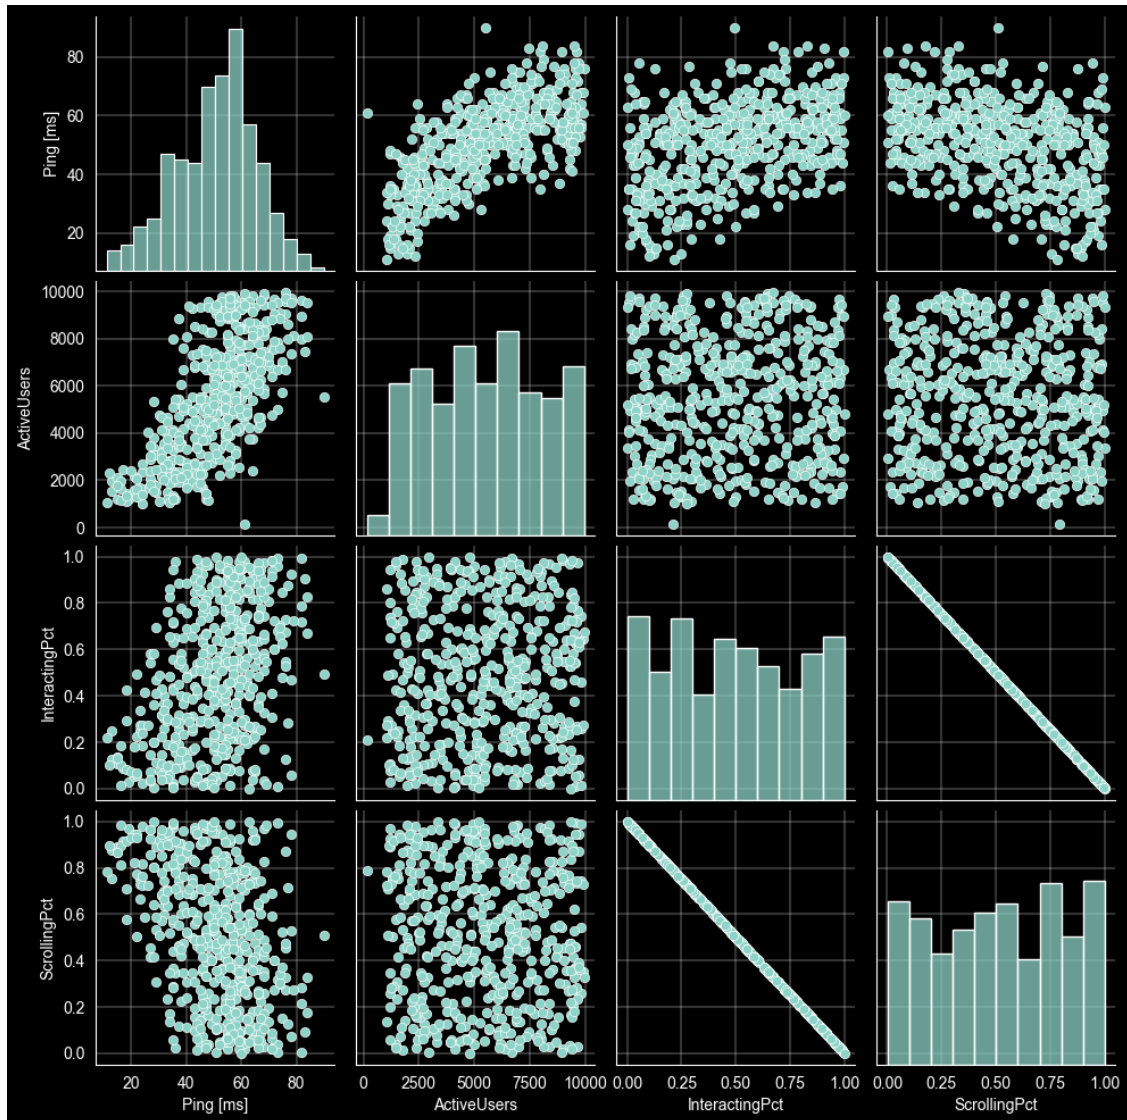
# Print data info to learn more about data
print(df_uloha_2.head())
print(df_uloha_2.describe())
```

	OSType	ActiveUsers	InteractingPct	ScrollingPct	Ping [ms]
0	iOS	4113	0.8283	0.1717	47
1	iOS	7549	0.3461	0.6539	46
2	Windows	8855	0.2178	0.7822	55
3	Android	8870	0.0794	0.9206	56
4	MacOS	9559	0.7282	0.2718	76

	ActiveUsers	InteractingPct	ScrollingPct	Ping [ms]
count	502.000000	502.000000	502.000000	502.000000
mean	5485.830677	0.488613	0.511387	50.545817
std	2548.935679	0.296000	0.296000	14.797937
min	153.000000	0.000500	0.001400	11.000000
25%	3357.500000	0.229300	0.257525	40.000000
50%	5456.000000	0.482950	0.517050	52.000000
75%	7461.500000	0.742475	0.770700	60.000000
max	9953.000000	0.998600	0.999500	90.000000

5.2.1 Visualize data using matrix plot

```
[149]: # Visualize data
ax = sns.pairplot(df_uloha_2[['Ping [ms]', 'ActiveUsers', 'InteractingPct',
↪ 'ScrollingPct']])
plt.show()
```



Based on the previous correlation matrix, we can see that there is a high correlation between InteractingPct and ScrollingPct. Therefore we can remove one of them.

I choose to remove **ScrollingPct**

```
[150]: # Remove correlated parameters
X = pd.DataFrame({
    'ActiveUsers': df_uloha_2.loc[:, 'ActiveUsers'],
    'InteractingPct': df_uloha_2.loc[:, 'InteractingPct'],
    'ScrollingPct': df_uloha_2.loc[:, 'ScrollingPct'],
})

# Standardize
```

```

X['ActiveUsers'] = (X['ActiveUsers'] - X['ActiveUsers'].mean()) /  $\sqrt{\text{X['ActiveUsers'].std()}}$ 
X['InteractingPct'] = (X['InteractingPct'] - X['InteractingPct'].mean()) /  $\sqrt{\text{X['InteractingPct'].std()}}$ 

# OS types Instead True/False values, we can use 1/0
X['Windows'] = df_uloha_2['OSType'].apply(lambda x: 1 if x == 'Windows' else 0)
X['iOS'] = df_uloha_2['OSType'].apply(lambda x: 1 if x == 'iOS' else 0)
X['MacOS'] = df_uloha_2['OSType'].apply(lambda x: 1 if x == 'MacOS' else 0)
X['Android'] = df_uloha_2['OSType'].apply(lambda x: 1 if x == 'Android' else 0)
correlation_matrix = np.corrcoef(X.values.T)
corr_params = np.abs(correlation_matrix) > 0.7
# Print all correlated parameters that are not on the main diagonal and those
  only above main diagonal
print("Correlated parameters:")
for i in range(corr_params.shape[0]):
    for j in range(corr_params.shape[1]):
        if i != j and i < j and corr_params[i, j]:
            print(f"{X.columns[i]} - {X.columns[j]}")
            print(f"Removing {X.columns[j]}")
            X = X.drop(X.columns[j], axis=1)
X.head()

```

Correlated parameters:
 InteractingPct - ScrollingPct
 Removing ScrollingPct

```

[150]:
ActiveUsers  InteractingPct  Windows  iOS  MacOS  Android
0      -0.538590         1.147592        0     1      0         0
1       0.809424        -0.481464        0     1      0         0
2       1.321795        -0.914910        1     0      0         0
3       1.327679        -1.382478        0     0      0         1
4       1.597988         0.809416        0     0      1         0

```

```

[151]: # Polynomial degree
degree = 2

# Use PolynomialFeatures
poly = PolynomialFeatures(degree=degree, include_bias=True)
poly_features = poly.fit_transform(X)

# Create a new dataframe with the polynomial features and original column names
poly_X = pd.DataFrame(poly_features, columns=poly.get_feature_names_out(X.
  columns))

# Rename 1 to const
poly_X.rename(columns={'1': 'const'}, inplace=True)

```

```
# poly_X
```

```
[152]: def get_column_to_remove(model):  
    """  
    Firstly get all quadratic columns ending with ^2, then remove interaction_  
    ↪ terms and after all linear terms  
    :param model:  
    :return:  
    """  
    pvalues = model.pvalues  
  
    # Find all columns with p-value > 0.05 and nan  
    pvalues = pvalues[(pvalues > 0.05) | (pvalues.isna())]  
    pvalues = pvalues.drop('const') if 'const' in pvalues else pvalues  
  
    # Check if there is any quadratic term  
    quadratic_terms = [i for i in pvalues.index if i.endswith('^2')]  
  
    # Check if there is any interaction term  
    interaction_terms = [i for i in pvalues.index if ' ' in i]  
  
    # Check if there is any linear term  
    linear_terms = [i for i in pvalues.index if i not in quadratic_terms and i_  
    ↪ not in interaction_terms]  
  
    # Find nan values  
    nan_values = [i for i in pvalues.index if  
                   i not in quadratic_terms and i not in interaction_terms and i_  
    ↪ not in linear_terms]  
  
    if len(quadratic_terms) > 0:  
        return quadratic_terms[0]  
    elif len(interaction_terms) > 0:  
        return interaction_terms[0]  
    elif len(linear_terms) > 0:  
        return linear_terms[0]  
    elif len(nan_values) > 0:  
        return nan_values[0]  
    else:  
        return None
```

```
[153]: # Train  
y = df_uloha_2['Ping [ms]']  
model = sm.OLS(endog=y, exog=poly_X).fit()  
  
# Remove from poly_X the values that has p-value >= 0.05  
while remove_col := get_column_to_remove(model):
```

```

print(f"Removing {remove_col}")
poly_X = poly_X.drop(remove_col, axis=1) # remove column from X
model = sm.OLS(endog=y, exog=poly_X).fit() # fit model again

# Print summary
print(model.summary())
write_to_file('tmp/out/model_summary_pvalue.txt', model.summary().as_text())

```

Removing InteractingPct^2
 Removing ActiveUsers Windows
 Removing ActiveUsers iOS
 Removing InteractingPct Android
 Removing InteractingPct Windows
 Removing InteractingPct iOS
 Removing InteractingPct MacOS
 Removing Windows iOS
 Removing Windows MacOS
 Removing Windows Android
 Removing iOS MacOS
 Removing iOS Android
 Removing MacOS Android

OLS Regression Results

```

=====
Dep. Variable:          Ping [ms]      R-squared:                0.843
Model:                  OLS            Adj. R-squared:          0.840
Method:                 Least Squares   F-statistic:              293.7
Date:                  Sat, 16 Dec 2023 Prob (F-statistic):       1.62e-191
Time:                  22:24:53         Log-Likelihood:           -1599.6
No. Observations:      502             AIC:                     3219.
Df Residuals:          492             BIC:                     3261.
Df Model:               9
Covariance Type:       nonrobust
=====

```

```

=====

```

		coef	std err	t	P> t

const		35.2506	0.258	136.475	0.000
34.743	35.758				
ActiveUsers		7.7862	0.367	21.210	0.000
7.065	8.507				
InteractingPct		5.0493	0.266	18.977	0.000
4.527	5.572				
Windows		9.8027	0.233	42.041	0.000
9.345	10.261				
iOS		5.0093	0.246	20.331	0.000
4.525	5.493				

MacOS		12.5724	0.229	54.900	0.000
12.122	13.022				
Android		7.8661	0.249	31.600	0.000
7.377	8.355				
ActiveUsers^2		-2.6838	0.285	-9.432	0.000
-3.243	-2.125				
ActiveUsers InteractingPct		-2.3187	0.269	-8.621	0.000
-2.847	-1.790				
ActiveUsers MacOS		5.8465	0.633	9.232	0.000
4.602	7.091				
ActiveUsers Android		2.2256	0.690	3.225	0.001
0.870	3.582				
Windows^2		9.8027	0.233	42.041	0.000
9.345	10.261				
iOS^2		5.0093	0.246	20.331	0.000
4.525	5.493				
MacOS^2		12.5724	0.229	54.900	0.000
12.122	13.022				
Android^2		7.8661	0.249	31.600	0.000
7.377	8.355				

Omnibus:	228.381	Durbin-Watson:	1.925
Prob(Omnibus):	0.000	Jarque-Bera (JB):	3196.157
Skew:	1.598	Prob(JB):	0.00
Kurtosis:	14.941	Cond. No.	6.71e+16

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The smallest eigenvalue is 3.27e-31. This might indicate that there are strong multicollinearity problems or that the design matrix is singular.

```
[154]: def get_column_to_remove_vif(df):
        """
        Firstly get all quadratic columns ending with ^2, then remove interaction
        ↪ terms and after all linear terms
        :param model:
        :return:
        """
        # Calculate vif
        vif = pd.Series([variance_inflation_factor(df.values, i) for i in range(df.
        ↪ shape[1])], index=df.columns)

        # Remove all values above 5 (infinite included), const can not be removed
        vif = vif[vif > 5]
```

```

# Don't remove const
vif = vif.drop('const') if 'const' in vif else vif

# Check if there is any quadratic term
quadratic_terms = [i for i in vif.index if i.endswith('^2')]

# Check if there is any interaction term
interaction_terms = [i for i in vif.index if ' ' in i]

# Check if there is any linear term
linear_terms = [i for i in vif.index if i not in quadratic_terms and i not
↳ in interaction_terms]

# Find nan values
nan_values = [i for i in vif.index if
               i not in quadratic_terms and i not in interaction_terms and i
↳ not in linear_terms]

if len(quadratic_terms) > 0:
    return quadratic_terms[0]
elif len(interaction_terms) > 0:
    return interaction_terms[0]
elif len(linear_terms) > 0:
    return linear_terms[0]
elif len(nan_values) > 0:
    return nan_values[0]
else:
    return None

```

```

[155]: import warnings

# Ignore warnings, because of division by zero when calculating vif
warnings.simplefilter("ignore", category=RuntimeWarning)

# Remove all parameters that has vif >= 5 (infinite included), const can not be
↳ removed
while remove_col := get_column_to_remove_vif(poly_X):
    print(f"Removing {remove_col}")
    poly_X = poly_X.drop(remove_col, axis=1)
    model = sm.OLS(endog=y, exog=poly_X).fit()

# Reset warnings to default
warnings.resetwarnings()

# Print summary
print(model.summary())
write_to_file('tmp/out/model_summary_vif.txt', model.summary().as_text())

```

```
# Calculate VIF
vif = pd.Series([variance_inflation_factor(poly_X.values, i) for i in
    range(poly_X.shape[1])], index=poly_X.columns)
vif
```

Removing Windows²
 Removing iOS²
 Removing MacOS²
 Removing Android²
 Removing Windows

OLS Regression Results

```
=====
Dep. Variable:          Ping [ms]      R-squared:                0.843
Model:                  OLS            Adj. R-squared:           0.840
Method:                 Least Squares   F-statistic:              293.7
Date:                  Sat, 16 Dec 2023 Prob (F-statistic):       1.62e-191
Time:                  22:24:53         Log-Likelihood:           -1599.6
No. Observations:      502             AIC:                     3219.
Df Residuals:          492             BIC:                     3261.
Df Model:               9
Covariance Type:        nonrobust
=====
```

```
=====
                                coef    std err          t      P>|t|
-----
[0.025    0.975]
-----
const                54.8560    0.591    92.857    0.000
53.695    56.017
ActiveUsers           7.7862    0.367    21.210    0.000
7.065    8.507
InteractingPct        5.0493    0.266    18.977    0.000
4.527    5.572
iOS                 -9.5869    0.749   -12.804    0.000
-11.058   -8.116
MacOS                5.5393    0.720     7.696    0.000
4.125    6.954
Android             -3.8732    0.761    -5.088    0.000
-5.369   -2.377
ActiveUsers2        -2.6838    0.285    -9.432    0.000
-3.243   -2.125
ActiveUsers InteractingPct -2.3187    0.269    -8.621    0.000
-2.847   -1.790
ActiveUsers MacOS      5.8465    0.633     9.232    0.000
4.602    7.091
ActiveUsers Android    2.2256    0.690     3.225    0.001
0.870    3.582
```



```
=====
Omnibus:                228.381    Durbin-Watson:                1.925
Prob(Omnibus):           0.000    Jarque-Bera (JB):           3196.157
Skew:                    1.598    Prob(JB):                    0.00
Kurtosis:                14.941    Cond. No.                    7.07
=====
```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

```
[155]: const                5.006291
      ActiveUsers           1.929304
      InteractingPct        1.013519
      iOS                   1.446080
      MacOS                 1.481309
      Android               1.440922
      ActiveUsers^2         1.013961
      ActiveUsers InteractingPct 1.016595
      ActiveUsers MacOS     1.527704
      ActiveUsers Android   1.416742
      dtype: float64
```

5.2.2 ULOHA 2.1.1 Zapište rovnici Vašeho finálního modelu.

```
[156]: # Print equation
model_params = model.params.drop('const')
equation = f"ping = \n{model.params['const']:.5f}\n"
for k, v in model_params.items():
    equation += f"+ {v:.5f} * {k}\n"
print(equation)
```

```
ping =
54.85603
+ 7.78621 * ActiveUsers
+ 5.04932 * InteractingPct
+ -9.58693 * iOS
+ 5.53933 * MacOS
+ -3.87321 * Android
+ -2.68377 * ActiveUsers^2
+ -2.31866 * ActiveUsers InteractingPct
+ 5.84648 * ActiveUsers MacOS
+ 2.22559 * ActiveUsers Android
```

5.2.3 ULOHA 2.1.2 Diskutujte splnění předpokladů lineární regrese a základní regresní diagnostiky.

TODO

5.2.4 ULOHA 2.1.3 Pokud (až během regresního modelování) identifikujete některé „extrémně odlehle hodnoty“ můžete ty „nejodlehlejší“ hodnoty, po alespoň krátkém zdůvodnění, vyřadit.

TODO

```
[157]: def plot_diagnostic_subplots(model, title: str = 'Diagnostic Plots'):
    """
    Plot diagnostic subplots
    :param model:
    :param title:
    :return:
    """

    # Set up subplots
    fig, axes = plt.subplots(1, 4, figsize=(4*4, 4))

    # Set title for whole plots
    fig.suptitle(title, fontsize=16)

    # Residua vs. Fitted Values (diagnostic graph)
    sns.scatterplot(x=model.fittedvalues, y=model.resid, ax=axes[0])
    axes[0].set_title("Residua vs. Fitted Values")
    axes[0].set_xlabel("Fitted Values")
    axes[0].set_ylabel("Residua")

    # Normality reziduů (Q-Q plot)
    sm.qqplot(model.resid, line='s', ax=axes[1])
    axes[1].set_title("Q-Q plot reziduů")

    # Homoskedasticita (diagnostic graph)
    influence = model.get_influence()
    residuals_studentized = influence.resid_studentized_internal
    fitted_values = model.fittedvalues
    sns.scatterplot(x=fitted_values, y=np.sqrt(np.abs(residuals_studentized)),
    ↪ax=axes[2])
    axes[2].set_title("Square Root of Standardized Residuals vs. Fitted Values")
    axes[2].set_xlabel("Fitted Values")
    axes[2].set_ylabel("Square Root of Standardized Residuals")

    # Distribution of Residuals
    residuals = model.resid
    sns.histplot(residuals, kde=True, ax=axes[3])
    axes[3].set_title('Distribution of Residuals')
    axes[3].set_xlabel('Residuals')
    axes[3].set_ylabel('Count')

    # Adjust layout to prevent clipping of titles
    plt.tight_layout()
```

```


# Show the plots
_title = title.lower().replace(' ', '_')
plt.savefig(f"tmp/out/diagnostic_plots_{_title}.png")
plt.show()

```

```

[158]: plot_diagnostic_subplots(model, title='Diagnostic Plots: Before Removing_
↳Outliers')

```



```

[159]: # Fit an OLS model
ols_model = OLSInfluence(model)

```

```

[160]: # Standardized residuals
standardized_residuals = ols_model.resid_studentized_internal

# Identify outliers based on standardized residuals
outliers = np.abs(standardized_residuals) > 5
outliers[outliers == True].index

```

```

[160]: Index([255, 476], dtype='int64')

```

```

[161]: # Cook's distance
cooks_distance = ols_model.cooks_distance[0]

# Identify outliers based on Cook's distance
cooks_outliers = cooks_distance > 10 / poly_X.shape[0]
cooks_outliers[cooks_outliers == True].index

```

```

[161]: Index([255, 476], dtype='int64')

```

```

[162]: merged_outliers = list(set(outliers[outliers == True].index) |
↳set(cooks_outliers[cooks_outliers == True].index))
merged_outliers.sort()

# Remove outliers, if was not removed before
if len(poly_X) == len(X):

```

```
poly_X = poly_X.drop(merged_outliers, axis=0)
y = y.drop(merged_outliers, axis=0)
```

```
# poly_X
```

```
[163]: # Retrain model
model_without_outliers = sm.OLS(endog=y, exog=poly_X).fit()
print(model_without_outliers.summary())
write_to_file('tmp/out/model_summary_cook.txt', model_without_outliers.
    ↳summary().as_text())
```

OLS Regression Results

```
=====
Dep. Variable:          Ping [ms]      R-squared:                0.877
Model:                  OLS            Adj. R-squared:          0.875
Method:                 Least Squares  F-statistic:             388.1
Date:                  Sat, 16 Dec 2023 Prob (F-statistic):       1.43e-216
Time:                  22:24:55        Log-Likelihood:          -1529.5
No. Observations:      500            AIC:                    3079.
Df Residuals:          490            BIC:                    3121.
Df Model:              9
Covariance Type:       nonrobust
=====
```


```
=====
                                coef    std err          t      P>|t|
-----
[0.025    0.975]
-----
const                54.9364      0.525    104.738    0.000
53.906    55.967
ActiveUsers          7.7474      0.323    23.970    0.000
7.112     8.382
InteractingPct       5.1512      0.234    21.970    0.000
4.691     5.612
iOS                 -9.3373      0.660   -14.140    0.000
-10.635   -8.040
MacOS               5.3424      0.637     8.391    0.000
4.091     6.593
Android            -3.6638      0.671    -5.456    0.000
-4.983    -2.344
ActiveUsers^2       -2.9856      0.254   -11.764    0.000
-3.484    -2.487
ActiveUsers InteractingPct -2.5439      0.238   -10.693    0.000
-3.011    -2.076
ActiveUsers MacOS    6.7342      0.565    11.929    0.000
5.625     7.843
ActiveUsers Android  2.2951      0.608     3.777    0.000
1.101     3.489
```

```
=====
Omnibus:                0.799    Durbin-Watson:                1.981
Prob(Omnibus):          0.671    Jarque-Bera (JB):        0.865
Skew:                   0.002    Prob(JB):                0.649
Kurtosis:               2.796    Cond. No.                7.06
=====
```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

```
[164]: plot_diagnostic_subplots(model_without_outliers, title='Diagnostic Plots: After_
↳ Removing Outliers')
```



5.3 ULOHA 2.2. [1. body] - Pomocí Vašeho výsledného modelu identifikujte, pro které nastavení parametrů má odezva nejproblematictější hodnotu.

```
[165]: # Find max ping value
max_ping = model_without_outliers.predict().argmax()
max_ping
```

[165]: 10

5.4 ULOHA 2.3. [1. bod] - Odhadněte hodnotu odezvy uživatele s Windows, při průměrném nastavení ostatních parametrů a vypočtete konfidenční interval a predikční interval pro toto nastavení.

```
[166]: # Average values
mean_poly_X = poly_X.mean()

# Predict ping for user with Windows
predicted_ping = model_without_outliers.predict(mean_poly_X)
print(f"Predikovaná odezva uživatele s Windows: {predicted_ping.values[0]:.5f}")

# Calculate confidence interval
```

```

confidence_interval = model_without_outliers.get_prediction(mean_poly_X).
    ↪ conf_int()

# Calculate prediction interval
prediction_interval = model_without_outliers.get_prediction(mean_poly_X).
    ↪ conf_int(obs=True)

# Print confidence and prediction interval
print("\nKonfidenční interval:")
print(confidence_interval)
print("\nPredikční interval:")
print(prediction_interval)

```

Predikovaná odezva uživatele s Windows: 50.44600

Konfidenční interval:

[[49.98837369 50.90362631]]

Predikční interval:

[[40.20293692 60.68906308]]

5.5 ULOHA 2.4. [2. body] - Na základě jakýchkoli vypočtených charakteristik argumentujte, zdali je Váš model „vhodný“ pro další použití.