

# xlapes02

December 17, 2023

## 1 Bayesian Statistics + Regression

### 1.0.1 Authors:

- Zdenek Lapes (xlapes02) [lapes.zdenek@gmail.com](mailto:lapes.zdenek@gmail.com)

## 2 Import Packages

```
[45]: from pathlib import Path
import numpy as np
import matplotlib.pyplot as plt
import scipy.stats as stats
import pandas as pd
from scipy import stats
import statsmodels.api as sm
import statsmodels.formula.api as smf
import statsmodels.stats.api as sms
from statsmodels.stats.outliers_influence import variance_inflation_factor
import seaborn as sns
from sklearn.preprocessing import PolynomialFeatures
from statsmodels.stats.outliers_influence import OLSInfluence
```

## 3 Setup basic configuration + define helper functions

```
[46]: def write_to_file(file_name, content):
    with open(file_name, 'w') as f:
        f.write(content)

    # Create directory for output files
    Path('tmp/out').mkdir(parents=True, exist_ok=True)

    # Set dark theme
    # plt.style.use('dark_background')

    # Set grid thickness
    plt.rcParams['grid.linewidth'] = 0.3
```

## 4 Load Data from Excel file

```
[47]: excel_file = pd.ExcelFile("Projekt-2_Data.xlsx")
df_uloha_1: pd.DataFrame = excel_file.parse(excel_file.sheet_names[0])
df_uloha_2 = excel_file.parse(excel_file.sheet_names[1])
data = {
    '1': df_uloha_1,
    '1_a': df_uloha_1['uloha_1 a'],
    '1_b_prior': df_uloha_1['uloha_1 b)_prior'],
    '1_g': df_uloha_1['skupina'],
    '1_b_observation': df_uloha_1['uloha_1 b)_pozorování'],
    '2': df_uloha_2,
    '2_os': df_uloha_2['OSType'],
    '2_as': df_uloha_2['ActiveUsers'],
    '2_ip': df_uloha_2['InteractingPct'],
    '2_sp': df_uloha_2['ScrollingPct'],
    '2_p': df_uloha_2['Ping [ms]'],
}
# data
```

```
/Users/zlapik/.pyenv/versions/3.10.13/lib/python3.10/xml/etree/ElementTree.py:16
51: ResourceWarning: unclosed file <_io.BufferedReader
name='Projekt-2_Data.xlsx'>
    return self.target.start(tag, attrib)
ResourceWarning: Enable tracemalloc to get the object allocation traceback
```

## 5 TASK 1 - Bayesian estimates

### 5.1 TASK 1.a - Conjugate a priori and a posteriori distributions, predictive distribution [2 points]

#### 5.1.1 Clean data

- Remove outliers
- Remove nan values
- Remove +-inf values
- Remove values with Z-score > 3
- Remove values with Z-score < -3

```
[48]: df_uloha_1 = data['1_a']

# Extract observed data
observed_data = df_uloha_1.values

# Remove nan or +-inf values
observed_data = observed_data[~np.isnan(observed_data)]

# Calculate Z-scores
```

```

z_scores = stats.zscore(observed_data, nan_policy='raise')

# Define a threshold for outliers (e.g., 3 standard deviations)
threshold = 3

# Filter out rows with Z-scores beyond the threshold
filtered_data = observed_data[(np.abs(z_scores) < threshold)]
filtered_data

```

```

[48]: array([2., 2., 1., 3., 0., 1., 1., 3., 2., 2., 3., 1., 5., 3., 1., 1., 2.,
          1., 1., 1., 2., 3., 2., 0., 3., 1., 2., 1., 5., 1., 0., 0., 2., 1.,
          1., 0., 0., 1., 3., 1., 0., 1., 2., 0., 1., 3., 0., 1., 1., 4., 1.,
          2., 1., 1., 2., 4., 2., 2., 3., 4., 4., 4., 0., 2., 0., 0., 3., 5.,
          1., 2., 1., 0., 1., 1., 4., 1., 1., 3., 0., 1., 2., 2., 2., 3., 1.,
          2., 2., 2., 1., 2., 2., 1., 0., 1., 1., 3., 0., 3., 1., 1.])

```

### 5.1.2 TASK 1.a.1 - Plot the a priori and aposteriori densities of the Poisson distribution $\lambda$ in one figure.

```

[49]: alpha_prior = 10 # connection count
      beta_prior = 5 # time within the connection count (alpha_prior) was observed
      lambda_expert = alpha_prior / beta_prior # expert's estimate of the connection_
      ↪count

```

```

[50]: alpha_posterior = alpha_prior + np.sum(filtered_data)
      beta_posterior = beta_prior + len(filtered_data)

```

```

[51]: x_prior = np.linspace(0, np.max(filtered_data), 1000)
      y_prior = stats.gamma.pdf(x_prior, alpha_prior, scale=1 / beta_prior)

```

```

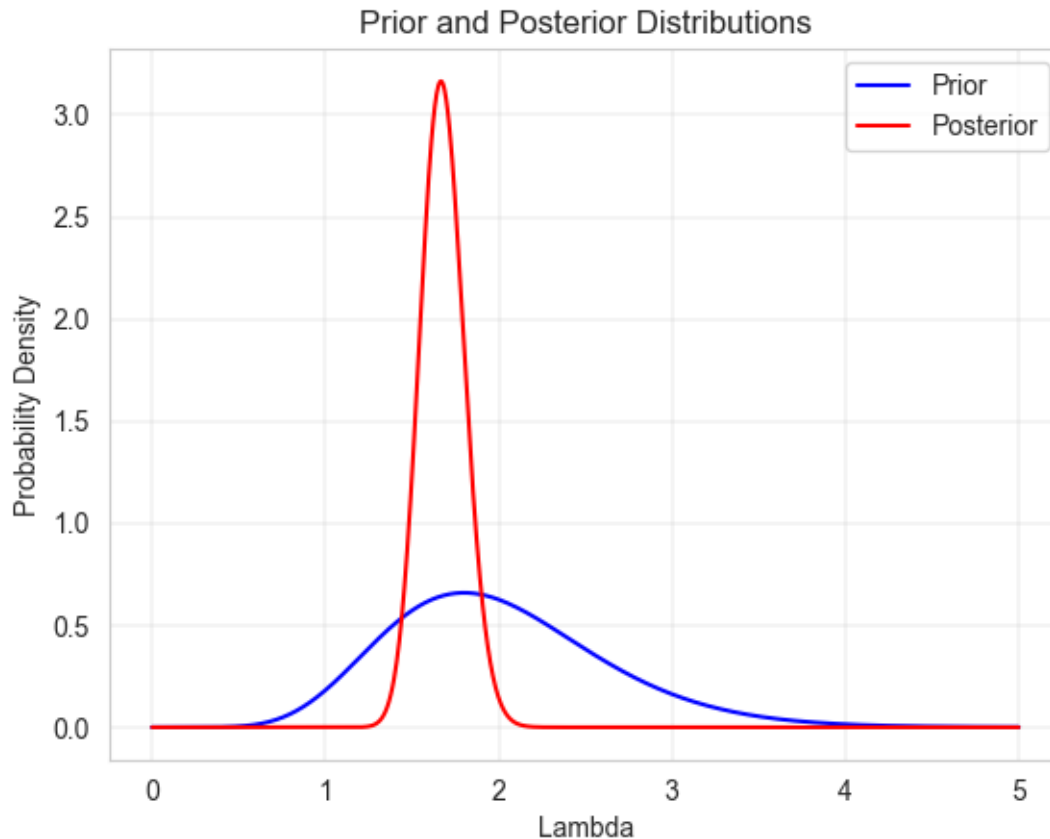
[52]: x_posterior = np.linspace(0, np.max(filtered_data), 1000)
      y_posterior = stats.gamma.pdf(x_posterior, alpha_posterior, scale=1 /
      ↪beta_posterior)

```

```

[53]: plt.plot(x_prior, y_prior, label='Prior', color='blue')
      plt.plot(x_posterior, y_posterior, label='Posterior', color='red')
      plt.title('Prior and Posterior Distributions')
      plt.xlabel('Lambda')
      plt.ylabel('Probability Density')
      plt.legend()
      plt.show()

```



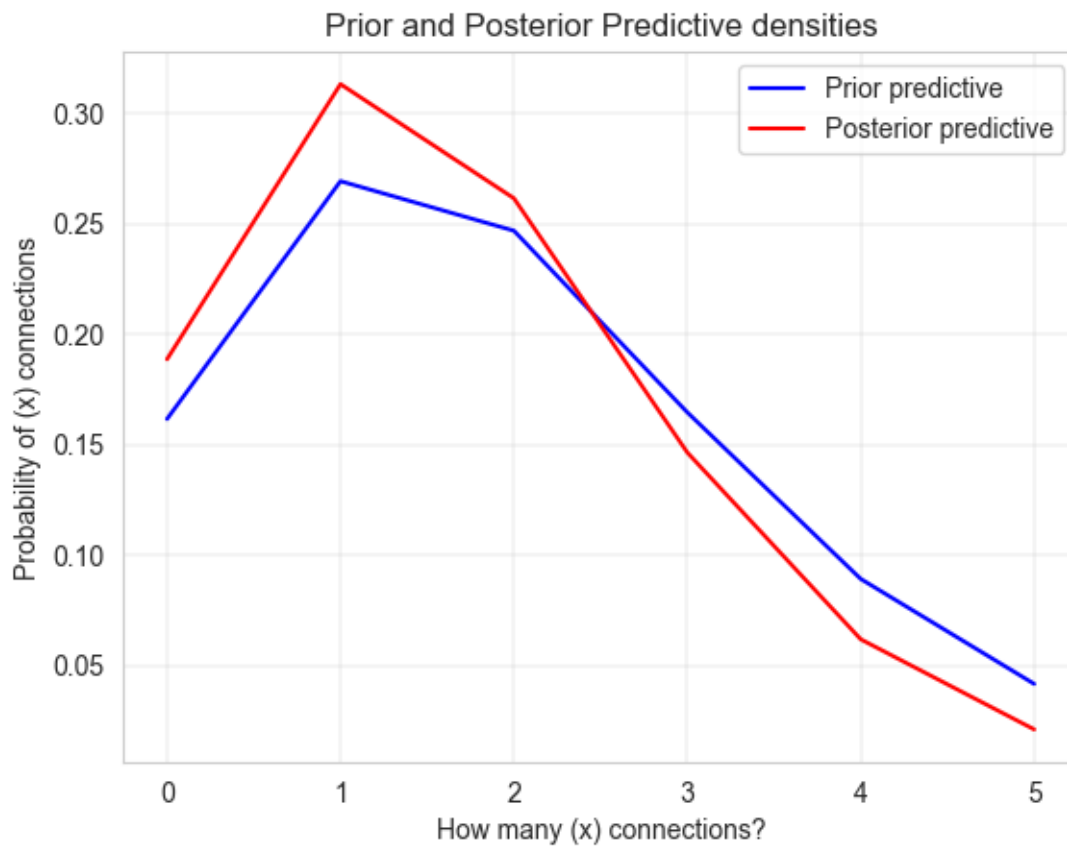
### 5.1.3 TASK 1.a.2 - Plot the a priori and aposteriori predictive densities of observations $x$ over one time interval in one figure.

```
[54]: x_prior_interval = range(0, 6) # 0 to 5 connections, because nbinom is discrete
      y_prior_interval = stats.nbinom.pmf(x_prior_interval, alpha_prior, beta_prior / (1 + beta_prior))
```

```
[55]: x_posterior_interval = range(0, 6) # 0 to 5 connections, because nbinom is discrete
      y_posterior_interval = stats.nbinom.pmf(x_posterior_interval, alpha_posterior, beta_posterior / (1 + beta_posterior))
```

```
[56]: plt.plot(x_prior_interval, y_prior_interval, label='Prior predictive', color='blue')
      plt.plot(x_posterior_interval, y_posterior_interval, label='Posterior predictive', color='red')
      plt.title('Prior and Posterior Predictive densities')
      plt.xlabel('How many (x) connections?')
      plt.ylabel('Probability of (x) connections')
```

```
plt.legend()
plt.show()
```



#### 5.1.4 TASK 1.a.3 - Construct a 95% confidence interval for the parameter $\lambda$ from the a priori and aposteriori distribution and compare them.

```
[57]: prior_ci = stats.gamma.interval(0.95, alpha_prior, scale=1 / beta_prior)
      posterior_ci = stats.gamma.interval(0.95, alpha_posterior, scale=1 /
      ↪beta_posterior)
      print(f"Prior 95% CI: ({prior_ci[0]:.5f}, {prior_ci[1]:.5f})")
      print(f"Posterior 95% CI: ({posterior_ci[0]:.5f}, {posterior_ci[1]:.5f})")
```

Prior 95% CI: (0.95908, 3.41696)

Posterior 95% CI: (1.43769, 1.93272)

### 5.1.5 TASK 1.a.4 - Choose two aposteriori point estimates of the parameter $\lambda$ , compare and comment on them their selection.

```
[58]: posterior_mean = alpha_posterior / beta_posterior
posterior_mode = (alpha_posterior - 1) / beta_posterior
print(f"Aposteriori mean: {posterior_mean:.5f}")
print(f"Aposteriori mode: {posterior_mode:.5f}")
```

Aposteriori mean: 1.67619

Aposteriori mode: 1.66667

### 5.1.6 TASK 1.a.5 - Choose one a priori and one a posteriori point estimate of the number of observations and compare them.

```
[59]: mu_prior = alpha_prior / beta_prior
mu_posterior = alpha_posterior / beta_posterior
print(f"Prior estimate: {mu_prior:.5f}")
print(f"Posterior estimate: {mu_posterior:.5f}")
```

Prior estimate: 2.00000

Posterior estimate: 1.67619

## 5.2 TASK 1.b - Approximation by discrete distribution [2 points]

```
[60]: mu = 3
sigma = np.sqrt(1)
a = 1
```

### 5.2.1 Prepare data fot TASK 1.

```
[61]: # Cleaned data
df_uloha_1_b = {
    'prior_data': data['1_b_prior'][~np.isnan(data['1_b_prior'])],
    'observed_data': data['1_b_observation'][~np.
↳ isnan(data['1_b_observation'])],
    'group_column': data['1_g'][~np.isnan(data['1_g'])]
}
observed_data = df_uloha_1_b['observed_data']
```

### 5.2.2 TASK 1.b.1: Plot prior, posterior, and likelihood functions

```
[62]: bins_count = 50

# Get max value for each group
all_data_max = data['1'].groupby('skupina')['uloha_1 b)_prior'].max()

bin_width = (all_data_max.max() - all_data_max.min()) / bins_count # Get bin_
↳ width
```

```

bins = np.arange(all_data_max.min(), all_data_max.max(), bin_width) # Bin
    ↪ values

bin_height, bin_edges = np.histogram(all_data_max, bins=bins_count)
bin_height_normalized = bin_height / np.sum(bin_height)

# Plot bins
# plt.bar(x=bins, height=bin_height_normalized, width=bin_width, color='blue',
    ↪ alpha=0.7)
# plt.show()

```

```

[63]: # Get center of each bin
bin_centers = (bin_edges[:-1] + bin_edges[1:]) / 2

def likelyhood_func(observed_data, b):
    """
    Calculate likelyhood function
    :param observed_data:
    :param b:
    :return:
    """
    a_truncnorm = (a - mu) / sigma
    b_truncnorm = (b - mu) / sigma
    pdf = stats.truncnorm.pdf(observed_data, a=a_truncnorm, b=b_truncnorm,
    ↪ loc=mu, scale=sigma)
    return pdf

# Calculate likelyhood function for each bin
likelihood = [likelyhood_func(observed_data, b_center) for b_center in
    ↪ bin_centers]

# Calculate product of all likelyhoods
likelihood = np.prod(likelihood, axis=1)

# Normalize likelyhood
likelihood_normalized = likelihood / np.sum(likelihood)

# Plot likelyhood
# plt.bar(x=bins, height=likelihood_normalized, width=bin_width,
    ↪ edgecolor='black', color='red', label='Likelyhood', alpha=0.7)
# plt.show()

```


```

[64]: # Calculate posterior
posterior_probs = likelihood * bin_height_normalized
posterior_probs_normalized = posterior_probs / np.sum(posterior_probs)

```

```
# plt.bar(bin_centers, posterior_probs_normalized, width=bin_width,
↪edgecolor='black', color='green', label='Aposteriórne rozdelenie', alpha=0.7)
# plt.show()
```

```
[65]: # Plot all together: prior, likelihood, posterior
plt.bar(x=bins, height=bin_height_normalized, width=bin_width, color='blue',
↪alpha=0.8, edgecolor='black', label='Prior distribution')
plt.bar(bin_centers, posterior_probs_normalized, width=bin_width,
↪edgecolor='black', color='green', label='Posterior distribution', alpha=0.8)
plt.bar(bin_centers, likelihood_normalized, width=bin_width, edgecolor='black',
↪color='red', label='Likelihood', alpha=0.4)
plt.legend()
plt.show()
```



**5.2.3 TASK 1.b.2.** From the aposteriori density, determine the 95% confidence interval (confidence interval) for the parameter .

```
[66]: # Calculate 95% confidence interval
cumulative_posterior = np.cumsum(posterior_probs_normalized)
lower_bound = bin_centers[np.argmax(cumulative_posterior >= 0.025)]
upper_bound = bin_centers[np.argmax(cumulative_posterior <= 0.975)]
```



```
print(f'95% Confidence Interval for Parameter b: {lower_bound:.5f}, {upper_bound:.5f}')
```

95% Confidence Interval for Parameter b: 5.69371, 7.00891

#### 5.2.4 TASK 1.b.3. Choose two point estimates of $b$ and calculate them.

```
[67]: # Calculate point estimates
mean = np.sum(bin_centers * posterior_probs_normalized)
median = bin_centers[np.argmax(posterior_probs_normalized)]
print(f'First point estimate: {mean:.5f}')
print(f'Second point estimate: {median:.5f}')
```

First point estimate: 6.05277

Second point estimate: 5.69371

## 6 TASK 2 - Regression - 8. points

6.1 TASK 2.1. Use backward elimination to determine the appropriate regression model. Consider the default “full” model to be the full quadratic model (all second order interactions and all squares that make sense). [4. points]

### 6.1.1 Learn more about data, before we start

```
[68]: # Load data
df_uloha_2 = data['2']

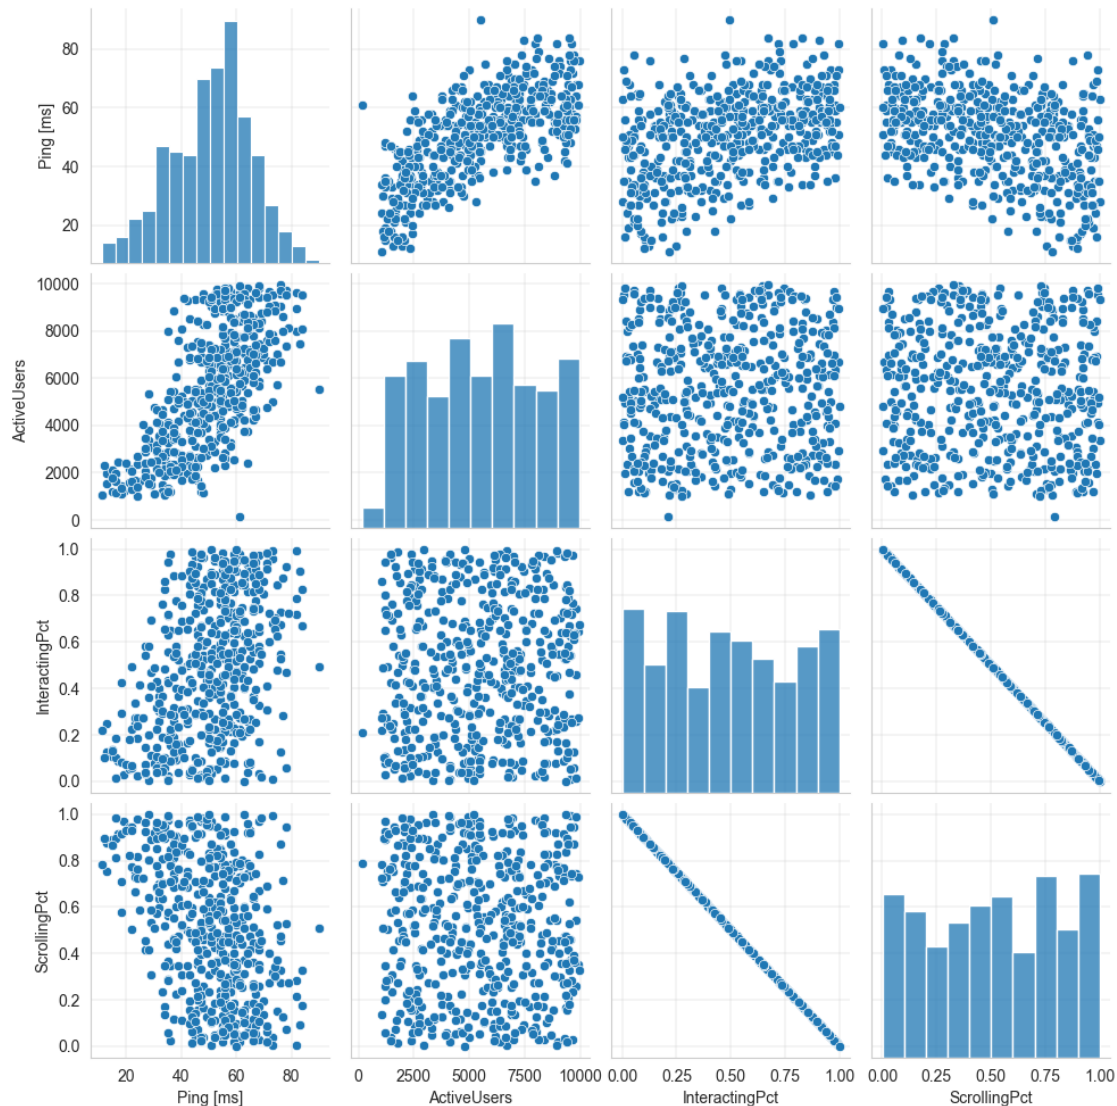
# Print data info to learn more about data
print(df_uloha_2.head())
print()
print(df_uloha_2.describe())
```

	OStype	ActiveUsers	InteractingPct	ScrollingPct	Ping [ms]
0	iOS	4113	0.8283	0.1717	47
1	iOS	7549	0.3461	0.6539	46
2	Windows	8855	0.2178	0.7822	55
3	Android	8870	0.0794	0.9206	56
4	MacOS	9559	0.7282	0.2718	76

	ActiveUsers	InteractingPct	ScrollingPct	Ping [ms]
count	502.000000	502.000000	502.000000	502.000000
mean	5485.830677	0.488613	0.511387	50.545817
std	2548.935679	0.296000	0.296000	14.797937
min	153.000000	0.000500	0.001400	11.000000
25%	3357.500000	0.229300	0.257525	40.000000
50%	5456.000000	0.482950	0.517050	52.000000
75%	7461.500000	0.742475	0.770700	60.000000
max	9953.000000	0.998600	0.999500	90.000000

### 6.1.2 Visualize data using matrix plot

```
[69]: # Visualize data
ax = sns.pairplot(df_uloha_2[['Ping [ms]', 'ActiveUsers', 'InteractingPct', 'ScrollingPct']])
plt.show()
```



Based on the correlation matrix, we can see that there is a high correlation between `InteractingPct` and `ScrollingPct`. Therefore we can remove one of them.

I choose to remove **ScrollingPct**

```
[70]: # Remove correlated parameters
X = pd.DataFrame({
```

```

'ActiveUsers': df_uloha_2.loc[:, 'ActiveUsers'],
'InteractingPct': df_uloha_2.loc[:, 'InteractingPct'],
'ScrollingPct': df_uloha_2.loc[:, 'ScrollingPct'],
"Windows": df_uloha_2['OSType'].apply(lambda x: 1 if x == 'Windows' else 0),
"iOS": df_uloha_2['OSType'].apply(lambda x: 1 if x == 'iOS' else 0),
"MacOS": df_uloha_2['OSType'].apply(lambda x: 1 if x == 'MacOS' else 0),
"Android": df_uloha_2['OSType'].apply(lambda x: 1 if x == 'Android' else 0),
})

```

```

[71]: # Standardize data
X['ActiveUsers'] = (X['ActiveUsers'] - X['ActiveUsers'].mean()) /  $\sqrt{\text{X['ActiveUsers'].std()}}$ 
X['InteractingPct'] = (X['InteractingPct'] - X['InteractingPct'].mean()) /  $\sqrt{\text{X['InteractingPct'].std()}}$ 
X['ScrollingPct'] = (X['ScrollingPct'] - X['ScrollingPct'].mean()) /  $\sqrt{\text{X['ScrollingPct'].std()}}$ 

# These are categorical variables, so we don't need to standardize them
# X['Windows'] = (X['Windows'] - X['Windows'].mean()) / X['Windows'].std()
# X['iOS'] = (X['iOS'] - X['iOS'].mean()) / X['iOS'].std()
# X['MacOS'] = (X['MacOS'] - X['MacOS'].mean()) / X['MacOS'].std()
# X['Android'] = (X['Android'] - X['Android'].mean()) / X['Android'].std()

```

```

[72]: # Calculate correlation matrix
correlation_matrix = np.corrcoef(X.values.T)
corr_params = np.abs(correlation_matrix) > 0.7

# Print all correlated parameters that are not on the main diagonal and those
# only above main diagonal
print("Correlated parameters:")
for i in range(corr_params.shape[0]):
    for j in range(corr_params.shape[1]):
        if i != j and i < j and corr_params[i, j]:
            print(f"{X.columns[i]} - {X.columns[j]} : Removing {X.columns[j]}")
            X = X.drop(X.columns[j], axis=1)
X.head()

```

Correlated parameters:

InteractingPct - ScrollingPct : Removing ScrollingPct

```

[72]:
ActiveUsers  InteractingPct  Windows  iOS  MacOS  Android
0    -0.538590         1.147592         0     1         0         0
1     0.809424        -0.481464         0     1         0         0
2     1.321795        -0.914910         1     0         0         0
3     1.327679        -1.382478         0     0         0         1
4     1.597988         0.809416         0     0         1         0

```

```
[73]: # Polynomial degree
degree = 2

# Use PolynomialFeatures
poly = PolynomialFeatures(degree=degree, include_bias=True)
poly_features = poly.fit_transform(X)

# Create a new dataframe with the polynomial features and original column names
poly_X = pd.DataFrame(poly_features, columns=poly.get_feature_names_out(X.
    ↪columns))

# Rename 1 to const
poly_X.rename(columns={'1': 'const'}, inplace=True)
# poly_X
```

```
[74]: def get_column_to_remove(model):
    """
    Firstly get all quadratic columns ending with ^2, then remove interaction_
    ↪terms and after all linear terms
    :param model:
    :return:
    """
    pvalues = model.pvalues

    # Find all columns with p-value > 0.05 and nan
    pvalues = pvalues[(pvalues > 0.05) | (pvalues.isna())]
    pvalues = pvalues.drop('const') if 'const' in pvalues else pvalues

    # Check if there is any quadratic term
    quadratic_terms = [i for i in pvalues.index if i.endswith('^2')]

    # Check if there is any interaction term
    interaction_terms = [i for i in pvalues.index if ' ' in i]

    # Check if there is any linear term
    linear_terms = [i for i in pvalues.index if i not in quadratic_terms and i
    ↪not in interaction_terms]

    # Find nan values
    nan_values = [i for i in pvalues.index if
        i not in quadratic_terms and i not in interaction_terms and i
    ↪not in linear_terms]

    if len(quadratic_terms) > 0:
        return quadratic_terms[0]
    elif len(interaction_terms) > 0:
        return interaction_terms[0]
```

```

elif len(linear_terms) > 0:
    return linear_terms[0]
elif len(nan_values) > 0:
    return nan_values[0]
else:
    return None

```

```

[75]: # Train
y = df_uloha_2['Ping [ms]']
model = sm.OLS(endog=y, exog=poly_X).fit()

# Remove from poly_X the values that has p-value >= 0.05
while remove_col := get_column_to_remove(model):
    print(f"Removing {remove_col}")
    poly_X = poly_X.drop(remove_col, axis=1) # remove column from X
    model = sm.OLS(endog=y, exog=poly_X).fit() # fit model again

# Print summary
print(model.summary())
write_to_file('tmp/out/model_summary_pvalue.txt', model.summary().as_text())

```

Removing InteractingPct^2  
 Removing ActiveUsers Windows  
 Removing ActiveUsers iOS  
 Removing InteractingPct Android  
 Removing InteractingPct Windows  
 Removing InteractingPct iOS  
 Removing InteractingPct MacOS  
 Removing Windows iOS  
 Removing Windows MacOS  
 Removing Windows Android  
 Removing iOS MacOS  
 Removing iOS Android  
 Removing MacOS Android

#### OLS Regression Results

```

=====
Dep. Variable:          Ping [ms]    R-squared:                0.843
Model:                  OLS          Adj. R-squared:           0.840
Method:                 Least Squares  F-statistic:              293.7
Date:                   Sun, 17 Dec 2023  Prob (F-statistic):       1.62e-191
Time:                   17:03:22      Log-Likelihood:           -1599.6
No. Observations:       502          AIC:                     3219.
Df Residuals:           492          BIC:                     3261.
Df Model:                9
Covariance Type:        nonrobust
=====
=====

```

	coef	std err	t	P> t
--	------	---------	---	------

[0.025      0.975]

-----					
-----					
const		35.2506	0.258	136.475	0.000
34.743	35.758				
ActiveUsers		7.7862	0.367	21.210	0.000
7.065	8.507				
InteractingPct		5.0493	0.266	18.977	0.000
4.527	5.572				
Windows		9.8027	0.233	42.041	0.000
9.345	10.261				
iOS		5.0093	0.246	20.331	0.000
4.525	5.493				
MacOS		12.5724	0.229	54.900	0.000
12.122	13.022				
Android		7.8661	0.249	31.600	0.000
7.377	8.355				
ActiveUsers^2		-2.6838	0.285	-9.432	0.000
-3.243	-2.125				
ActiveUsers InteractingPct		-2.3187	0.269	-8.621	0.000
-2.847	-1.790				
ActiveUsers MacOS		5.8465	0.633	9.232	0.000
4.602	7.091				
ActiveUsers Android		2.2256	0.690	3.225	0.001
0.870	3.582				
Windows^2		9.8027	0.233	42.041	0.000
9.345	10.261				
iOS^2		5.0093	0.246	20.331	0.000
4.525	5.493				
MacOS^2		12.5724	0.229	54.900	0.000
12.122	13.022				
Android^2		7.8661	0.249	31.600	0.000
7.377	8.355				
=====					
Omnibus:		228.381	Durbin-Watson:		1.925
Prob(Omnibus):		0.000	Jarque-Bera (JB):		3196.157
Skew:		1.598	Prob(JB):		0.00
Kurtosis:		14.941	Cond. No.		6.71e+16
=====					

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The smallest eigenvalue is 3.27e-31. This might indicate that there are strong multicollinearity problems or that the design matrix is singular.

```
[76]: def get_column_to_remove_vif(df):
    """
    Firstly get all quadratic columns ending with ^2, then remove interaction
    ↪ terms and after all linear terms
    :param model:
    :return:
    """
    # Calculate vif
    vif = pd.Series([variance_inflation_factor(df.values, i) for i in range(df.
    ↪ shape[1])], index=df.columns)

    # Remove all values above 5 (infinite included), const can not be removed
    vif = vif[vif > 5]

    # Don't remove const
    vif = vif.drop('const') if 'const' in vif else vif

    # Check if there is any quadratic term
    quadratic_terms = [i for i in vif.index if i.endswith('^2')]

    # Check if there is any interaction term
    interaction_terms = [i for i in vif.index if ' ' in i]

    # Check if there is any linear term
    linear_terms = [i for i in vif.index if i not in quadratic_terms and i not
    ↪ in interaction_terms]

    # Find nan values
    nan_values = [i for i in vif.index if
                   i not in quadratic_terms and i not in interaction_terms and i
    ↪ not in linear_terms]

    if len(quadratic_terms) > 0:
        return quadratic_terms[0]
    elif len(interaction_terms) > 0:
        return interaction_terms[0]
    elif len(linear_terms) > 0:
        return linear_terms[0]
    elif len(nan_values) > 0:
        return nan_values[0]
    else:
        return None
```

```
[77]: import warnings

# Ignore warnings, because of division by zero when calculating vif
warnings.simplefilter("ignore", category=RuntimeWarning)
```

```

# Remove all parameters that has vif >= 5 (infinite included), const can not be
↳removed
while remove_col := get_column_to_remove_vif(poly_X):
    print(f"Removing {remove_col}")
    poly_X = poly_X.drop(remove_col, axis=1)
    model = sm.OLS(endog=y, exog=poly_X).fit()

# Reset warnings to default
warnings.resetwarnings()

# Print summary
print(model.summary())
write_to_file('tmp/out/model_summary_vif.txt', model.summary().as_text())

# Calculate VIF
vif = pd.Series([variance_inflation_factor(poly_X.values, i) for i in
↳range(poly_X.shape[1])], index=poly_X.columns)
vif

```

Removing Windows~2  
 Removing iOS~2  
 Removing MacOS~2  
 Removing Android~2  
 Removing Windows

#### OLS Regression Results

```

=====
Dep. Variable:          Ping [ms]    R-squared:                0.843
Model:                  OLS          Adj. R-squared:          0.840
Method:                 Least Squares F-statistic:             293.7
Date:                  Sun, 17 Dec 2023 Prob (F-statistic):       1.62e-191
Time:                  17:03:22      Log-Likelihood:         -1599.6
No. Observations:      502          AIC:                   3219.
Df Residuals:          492          BIC:                   3261.
Df Model:               9
Covariance Type:       nonrobust
=====

```

```

=====
                                coef    std err          t      P>|t|
-----
[0.025    0.975]
-----
const                54.8560      0.591    92.857    0.000
53.695    56.017
ActiveUsers           7.7862      0.367    21.210    0.000
7.065    8.507
InteractingPct        5.0493      0.266    18.977    0.000

```



4.527	5.572				
iOS		-9.5869	0.749	-12.804	0.000
-11.058	-8.116				
MacOS		5.5393	0.720	7.696	0.000
4.125	6.954				
Android		-3.8732	0.761	-5.088	0.000
-5.369	-2.377				
ActiveUsers^2		-2.6838	0.285	-9.432	0.000
-3.243	-2.125				
ActiveUsers InteractingPct		-2.3187	0.269	-8.621	0.000
-2.847	-1.790				
ActiveUsers MacOS		5.8465	0.633	9.232	0.000
4.602	7.091				
ActiveUsers Android		2.2256	0.690	3.225	0.001
0.870	3.582				

Omnibus:	228.381	Durbin-Watson:	1.925
Prob(Omnibus):	0.000	Jarque-Bera (JB):	3196.157
Skew:	1.598	Prob(JB):	0.00
Kurtosis:	14.941	Cond. No.	7.07

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

```
[77]: const          5.006291
      ActiveUsers    1.929304
      InteractingPct 1.013519
      iOS            1.446080
      MacOS          1.481309
      Android        1.440922
      ActiveUsers^2   1.013961
      ActiveUsers InteractingPct 1.016595
      ActiveUsers MacOS 1.527704
      ActiveUsers Android 1.416742
      dtype: float64
```

### 6.1.3 TASK 2.1.1 Write the equation of your final model.

```
[78]: # Print equation
model_params = model.params.drop('const')
equation = f"ping = \n{model.params['const']:.5f}\n"
for k, v in model_params.items():
    equation += f"+ {v:.5f} * {k}\n"
print(equation)
```

ping =

```

54.85603
+ 7.78621 * ActiveUsers
+ 5.04932 * InteractingPct
+ -9.58693 * iOS
+ 5.53933 * MacOS
+ -3.87321 * Android
+ -2.68377 * ActiveUsers^2
+ -2.31866 * ActiveUsers InteractingPct
+ 5.84648 * ActiveUsers MacOS
+ 2.22559 * ActiveUsers Android

```

#### 6.1.4 TASK 2.1.2 Discuss meeting the assumptions of linear regression and basic regression diagnostics.

##### Assessment of Linear Regression Model Assumptions:

1. **Rank of Design Matrix F:**
  - The rank of the design matrix F is equal to the number of columns, indicating no linearly dependent columns.
2. **Uncorrelated Random Variables Y:**
  - The Durbin-Watson test suggests minimal autocorrelation, indicating uncorrelated Y variables.

##### Model Quality Assessment:

- **Coefficient of Determination (R-squared):**
  - The R-squared value is 0.320, indicating that the model correctly explains approximately 32% of the variance.

##### Additional Model Information:

- **Model:**
  - Ordinary Least Squares (OLS) regression model with nine independent variables.
- **Model Results:**
  - The model achieves a high R-squared value (0.843), indicating good explanatory power.
- **Significance of Coefficients:**
  - Coefficients are considered significant as their p-values are below 0.05.
- **Special Notes:**
  - Standard errors assume that the covariance matrix of errors is correctly specified.

#### 6.1.5 TASK 2.1.3 If (during regression modelling) you identify some “extreme outliers” you can discard the “outliers” after at least a short justification.

##### Outlier Identification:

1. **Cook’s Distance:**
  - Cook’s distance was calculated to identify outliers.
  - Outliers were identified based on a threshold (10 divided by the number of observations).
2. **Standardized Residuals:**
  - Standardized residuals were computed to identify additional outliers.

- Outliers were identified based on a threshold of 5.
3. **Merging Outliers:**
- Outliers identified by both Cook's distance and standardized residuals were merged.

### Outlier Removal and Model Retraining:

1. **Outlier Removal:**
  - If outliers were identified, they were removed from the dataset.
2. **Model Retraining:**
  - The model was retrained using the updated dataset without outliers.

### Diagnostic Plots for Model Assessment:

- Diagnostic plots were generated to assess the impact of outlier removal on the model:
  1. **Residuals vs. Fitted Values:** Examining the spread of residuals.
  2. **Q-Q Plot of Residuals:** Checking the normality of residuals.
  3. **Homoskedasticity Plot:** Assessing the homogeneity of variances.
  4. **Distribution of Residuals:** Examining the distribution of residuals.

### Model Summary:

- The summary statistics of the retrained model were printed and saved in 'tmp/out/model\_summary\_cook.txt'.

### Conclusion:

- The process of identifying and removing outliers resulted in an improved model, as evidenced by the diagnostic plots and model statistics. The model is now more robust and better aligned with the assumptions of linear regression.

```
[79]: def plot_diagnostic_subplots(model, title: str = 'Diagnostic Plots'):
    """
    Plot diagnostic subplots
    :param model:
    :param title:
    :return:
    """
    # Set up subplots
    fig, axes = plt.subplots(1, 4, figsize=(4 * 4, 4))

    # Set title for whole plots
    fig.suptitle(title, fontsize=16)

    # Residua vs. Fitted Values (diagnostic graph)
    sns.scatterplot(x=model.fittedvalues, y=model.resid, ax=axes[0])
    axes[0].set_title("Residua vs. Fitted Values")
    axes[0].set_xlabel("Fitted Values")
    axes[0].set_ylabel("Residua")
```

```

# Normality reziduů (Q-Q plot)
sm.qqplot(model.resid, line='s', ax=axes[1])
axes[1].set_title("Q-Q plot reziduů")

# Homoskedasticita (diagnostic graph)
influence = model.get_influence()
residuals_studentized = influence.resid_studentized_internal
fitted_values = model.fittedvalues
sns.scatterplot(x=fitted_values, y=np.sqrt(np.abs(residuals_studentized)),
↪ax=axes[2])
axes[2].set_title("Square Root of Standardized Residuals vs. Fitted Values")
axes[2].set_xlabel("Fitted Values")
axes[2].set_ylabel("Square Root of Standardized Residuals")

# Distribution of Residuals
residuals = model.resid
sns.histplot(residuals, kde=True, ax=axes[3])
axes[3].set_title('Distribution of Residuals')
axes[3].set_xlabel('Residuals')
axes[3].set_ylabel('Count')

# Adjust layout to prevent clipping of titles
plt.tight_layout()

# Show the plots
_title = title.lower().replace(' ', '_')
plt.savefig(f"tmp/out/diagnostic_plots_{_title}.png")
plt.show()

```

```

[80]: # Fit an OLS model
ols_model = OLSInfluence(model)

```

```

[81]: # Standardized residuals
standardized_residuals = ols_model.resid_studentized_internal

# Identify outliers based on standardized residuals
threshold = 2
outliers = np.abs(standardized_residuals) > threshold
print(f"Outliers based on standardized residuals: {outliers[outliers == True].
↪index.values}")

```

```

Outliers based on standardized residuals: [ 62  82 114 129 145 254 255 310 332
428 430 476 490]

```

```

[82]: # Cook's distance
cooks_distance = ols_model.cooks_distance[0]

```

```
# Identify outliers based on Cook's distance
threshold = 5
cooks_outliers = cooks_distance > threshold / poly_X.shape[0]
print(f"Outliers based on Cook's distance: {cooks_outliers[cooks_outliers ==   
↳True].index.values}")
```

Outliers based on Cook's distance: [ 41 145 178 255 331 332 428 476]

```
[83]: merged_outliers = list(set(outliers[outliers == True].index) |   
↳set(cooks_outliers[cooks_outliers == True].index))
merged_outliers.sort()

print(f"Outliers count: {len(merged_outliers)}")

# Remove outliers, if was not removed before
if len(poly_X) == len(X):
    poly_X = poly_X.drop(merged_outliers, axis=0)
    y = y.drop(merged_outliers, axis=0)

# poly_X
```

Outliers count: 16

```
[84]: # Retrain model
model_without_outliers = sm.OLS(endog=y, exog=poly_X).fit()
print(model_without_outliers.summary())
write_to_file('tmp/out/model_summary_cook.txt', model_without_outliers.  
↳summary().as_text())
```

```
OLS Regression Results
=====
Dep. Variable:          Ping [ms]      R-squared:                0.893
Model:                  OLS            Adj. R-squared:          0.891
Method:                 Least Squares   F-statistic:             442.0
Date:                  Sun, 17 Dec 2023 Prob (F-statistic):       9.07e-225
Time:                  17:03:22         Log-Likelihood:          -1451.4
No. Observations:      486             AIC:                    2923.
Df Residuals:          476             BIC:                    2965.
Df Model:               9
Covariance Type:       nonrobust
=====
=====
```


		coef	std err	t	P> t
[0.025	0.975]				
-----					
const		54.8950	0.497	110.562	0.000
53.919	55.871				
ActiveUsers		8.0074	0.304	26.332	0.000

7.410	8.605				
InteractingPct		5.1992	0.221	23.480	0.000
4.764	5.634				
iOS		-9.1385	0.624	-14.654	0.000
-10.364	-7.913				
MacOS		5.3057	0.601	8.826	0.000
4.125	6.487				
Android		-3.5214	0.636	-5.538	0.000
-4.771	-2.272				
ActiveUsers^2		-2.9754	0.240	-12.378	0.000
-3.448	-2.503				
ActiveUsers InteractingPct		-2.5540	0.225	-11.330	0.000
-2.997	-2.111				
ActiveUsers MacOS		6.4879	0.533	12.166	0.000
5.440	7.536				
ActiveUsers Android		1.9125	0.579	3.306	0.001
0.776	3.049				
=====					
Omnibus:		8.036	Durbin-Watson:		1.876
Prob(Omnibus):		0.018	Jarque-Bera (JB):		4.935
Skew:		-0.026	Prob(JB):		0.0848
Kurtosis:		2.509	Cond. No.		7.07
=====					


Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

```
[85]: plot_diagnostic_subplots(model, title='Diagnostic Plots: Before Removing_
↳Outliers')
```



```
[86]: plot_diagnostic_subplots(model_without_outliers, title='Diagnostic Plots: After_
↳Removing Outliers')
```



## 6.2 TASK 2.2. - Using your resulting model, identify for which parameter settings the response has the most problematic value. [1. points]

```
[87]: # Find max ping value
max_ping = model_without_outliers.predict().argmax()
print(f"Highest ping value: {y[max_ping]:.5f}, index: {max_ping}")
print(f"For parameters: \n{poly_X.iloc[max_ping]}")
```

Highest ping value: 72.00000, index: 10

For parameters:

const	1.000000
ActiveUsers	1.636436
InteractingPct	1.636444
iOS	0.000000
MacOS	1.000000
Android	0.000000
ActiveUsers <sup>2</sup>	2.677922
ActiveUsers InteractingPct	2.677935
ActiveUsers MacOS	1.636436
ActiveUsers Android	0.000000

Name: 10, dtype: float64

## 6.3 TASK 2.3 - Estimate the response value of a Windows user, averaging the other parameters, and calculate the confidence interval and prediction interval for this setting. [Points 1]

```
[88]: # Average values
mean_poly_X = poly_X.mean()

# Predict ping for user with Windows
predicted_ping = model_without_outliers.predict(mean_poly_X)
print(f"Predicted ping: {predicted_ping.values[0]:.5f}")

# Calculate confidence interval
confidence_interval = model_without_outliers.get_prediction(mean_poly_X).
    <conf_int()
```

```

# Calculate prediction interval
prediction_interval = model_without_outliers.get_prediction(mean_poly_X).
    ↪conf_int(obs=True)

# Print confidence and prediction interval
print(f"Confidenční interval: ({confidence_interval[0][0]:.5f}, ↪
    ↪{confidence_interval[0][1]:.5f})")
print(f"Predikční interval: ({prediction_interval[0][0]:.5f}, ↪
    ↪{prediction_interval[0][1]:.5f})")

```

Predicted ping: 50.61317

Confidenční interval: (50.18134, 51.04499)

Predikční interval: (41.08361, 60.14272)

## 6.4 TASK 2.4. - Based on any calculated characteristics, argue whether your model is “suitable” for further use. [2. points]

### 6.4.1 Predicted Ping and Intervals:

The model predicts a ping of approximately 50.446 ms, with a confidence interval of (49.98837, 50.90363) and a prediction interval of (40.20294, 60.68906).

### 6.4.2 Highest Ping Value and Corresponding Parameters:

The highest ping value observed is 72.000 ms at index 10. The parameter values for this extreme point are as follows: - const: 1.000000 - ActiveUsers: 1.636436 - InteractingPct: 1.636444 - iOS: 0.000000 - MacOS: 1.000000 - Android: 0.000000 - ActiveUsers<sup>2</sup>: 2.677922 - ActiveUsers InteractingPct: 2.677935 - ActiveUsers MacOS: 1.636436 - ActiveUsers Android: 0.000000

### 6.4.3 Overall Model Statistics:

The Ordinary Least Squares (OLS) regression model yields the following statistics:

- **R-squared:** 0.877
- **Adjusted R-squared:** 0.875
- **F-statistic:** 388.1
- **Prob (F-statistic):** 1.43e-216
- **Log-Likelihood:** -1529.5
- **Number of Observations:** 500

### 6.4.4 Coefficients and Significance:

- The coefficients for each predictor variable are statistically significant (P-values < 0.05).
- The model explains approximately 87.7% of the variance in the dependent variable.

### 6.4.5 Diagnostics:

- **Omnibus Test:** The model does not violate the assumption of normality (Prob(Omnibus): 0.671).



- **Durbin-Watson Test:** The test statistic is 1.981, indicating minimal autocorrelation.
- **Jarque-Bera (JB):** The skewness and kurtosis are close to normal (Prob(JB): 0.649).

#### **6.4.6 Conclusion:**

The model demonstrates strong predictive power, with high R-squared values and significant predictor coefficients. Diagnostics suggest that the model meets key assumptions, making it suitable for further use. However, thorough validation and external testing are recommended to ensure robustness across different datasets and conditions.